

stylgen_v0 Architecture Diagrams

1. High-Level System Overview

```
graph TB
    Client[Client/Frontend]

    subgraph "FastAPI Application"
        API[API Routes<br/>main.py]

        subgraph "Core Components"
            Pipeline[Pipeline<br/>Orchestrator]
            Embedder[Embedder<br/>Text→Vector]
            LLM[LLM Provider<br/>Ollama/Dummy]
            VStore[Vector Store<br/>Per-user Index]
            Persona[Persona Builder]
        end

        subgraph "Storage Layer"
            Memory[Memory Store<br/>In-memory]
        end
    end

    subgraph "External Services"
        Ollama[Ollama Server<br/>localhost:11434]
    end

    Client -->|HTTP/SSE| API
    API --> Pipeline
    Pipeline --> Embedder
    Pipeline --> LLM
    Pipeline --> VStore
    API --> Memory
    Persona --> Embedder
    LLM -->|API calls| Ollama
```

2. Persona Creation Flow

```
sequenceDiagram
    participant C as Client
    participant A as API
    participant P as Persona Builder
    participant E as Embedder
    participant V as Vector Store
    participant M as Memory Store

    C->>A: POST /persona<br/>{user_id, samples[], preferences}
    A->>E: embed(samples)
```

```

E-->A: embeddings[]

A->>V: replace(user_id, vec_items[])
Note over V: Clears old samples<br/>Stores new ones

A->>P: build_persona(samples, preferences)
P->>E: embed(samples)
E-->P: embeddings[]
P-->P: Calculate centroid<br/>Select exemplar_ids
P-->A: PersonaCard

A->>M: store.personas[user_id] = PersonaCard
A->>M: store.samples[user_id] = WritingSample[]

A-->>C: PersonaCreateResponse<br/>{user_id, num_samples, persona}

```

3. Generation Pipeline (Detailed)

flowchart TB

Start([Generation Request])

subgraph "1. Retrieval Phase"

```

    GetPersona[Load PersonaCard<br/>from Memory]
    CheckExemplars{Has<br/>exemplar_ids?}
    UseExemplars[Fetch exemplars<br/>by ID]
    QueryVectors[Query vectors with<br/>centroid embedding]
    GetSamples[Get top-k similar<br/>samples via cosine]

```

end

subgraph "2. Prompt Construction"

```

    BuildSystem[Build System Prompt<br/>- Tone descriptors<br/>- Formality
level<br/>- Emoji/hashtag rules<br/>- Banned phrases<br/>- Structure preference]
    BuildUser[Build User Prompt<br/>- Brief keywords<br/>- Goal & audience<br/>-
Length hint<br/>- CTA & link<br/>- Exemplar snippets]

```

end

subgraph "3. LLM Generation"

```

    SetTemp{Custom<br/>temperature?}
    UseCustom[Use provided<br/>temperature]
    Alternate[Alternate between<br/>0.6 and 0.8]
    CallLLM[Call LLM Provider<br/>with prompts]
    HandleError{LLM<br/>available?}
    UseDummy[Use Dummy<br/>Provider]

```

end

subgraph "4. Critique & Polish"

```

    RemoveBanned[Remove banned phrases<br/>via regex]
    NormalizeWS[Normalize whitespace<br/>preserve line breaks]
    CheckLength[Check length<br/>60-140% of hint]

```

end

```

subgraph "5. Scoring"
    StyleSim[Calculate style similarity<br/>cosine(text, centroid)]
    Novelty[Calculate novelty<br/>1 - max_sim(text, samples)]
    Structure[Check structure<br/>requirements]
    CompositeScore[Composite scoring<br/>style + novelty + constraints]
end

subgraph "6. Selection"
    SortVariants[Sort variants<br/>by score]
    SelectBest[Select top as<br/>'chosen']
    Package[Package response<br/>with generation_id]
end

Start --> GetPersona
GetPersona --> CheckExemplars
CheckExemplars -->|Yes| UseExemplars
CheckExemplars -->|No| QueryVectors
QueryVectors --> GetSamples
UseExemplars --> BuildSystem
GetSamples --> BuildSystem

BuildSystem --> BuildUser
BuildUser --> SetTemp
SetTemp -->|Yes| UseCustom
SetTemp -->|No| Alternate
UseCustom --> CallLLM
Alternate --> CallLLM
CallLLM --> HandleError
HandleError -->|Success| RemoveBanned
HandleError -->|Fail| UseDummy
UseDummy --> RemoveBanned

RemoveBanned --> NormalizeWS
NormalizeWS --> CheckLength
CheckLength --> StyleSim
StyleSim --> Novelty
Novelty --> Structure
Structure --> CompositeScore
CompositeScore --> SortVariants
SortVariants --> SelectBest
SelectBest --> Package
Package --> End([Return Response])

```

4. Streaming Generation Flow

```

sequenceDiagram
    participant C as Client
    participant A as API
    participant P as Pipeline

```

```

participant L as LLM Provider
participant O as Ollama

C->>A: POST /generate/stream
A->>P: prepare_generation_context()
P->>A: exemplars, system_prompt

A->>C: SSE: event: meta<br/>data: {exemplars, goal, keywords}

A->>L: stream_generate()
L->>O: POST /api/generate<br/>{stream: true}

loop Token Streaming
  O->>L: {response: "token"}
  L->>A: yield "token"
  A->>C: SSE: data: token
end

O->>L: {done: true, eval_count, eval_duration}
L->>A: Log throughput metrics
A->>C: SSE: event: done<br/>data: end

```

5. Data Model Relationships

```

erDiagram
    USER ||--|| PERSONA-CARD : has
    USER ||--o{ WRITING-SAMPLE : owns
    USER ||--o{ GENERATION-RECORD : creates
    PERSONA-CARD ||--o{ EXEMPLAR-ID : references
    GENERATION-RECORD ||--o{ GENERATION-VARIANT : contains

    PERSONA-CARD {
        string user_id PK
        preferences object
        list exemplar_ids
        list centroid
    }

    WRITING-SAMPLE {
        string id PK
        string user_id FK
        string text
        list embedding
    }

    GENERATION-RECORD {
        string generation_id PK
        string user_id FK
        timestamp created_at
        brief object
    }

```

```

    list variants
}

GENERATION-VARIANT {
    string text
    score object
    float style_similarity
    float novelty
    bool structure_ok
    bool length_ok
}

```

6. Embedding & Retrieval Flow

```

flowchart LR
    subgraph "Text Processing"
        Input[Raw Text]
        Token[Tokenize<br/>lowercase, split]
        Hash[Hash tokens<br/>BLAKE2b]
        Vector[Build sparse vector<br/>dim=384]
        Norm[L2 Normalize]
    end

    subgraph "Storage"
        VecItem[VecItem<br/>{id, text, vec}]
        UserIndex[(Per-user Index)]
    end

    subgraph "Retrieval"
        Query[Query Vector]
        Cosine[Cosine Similarity<br/> $\text{dot}(a,b) / ||a|| \cdot ||b||$ ]
        TopK[Select top-k<br/>by similarity]
    end

    Input --> Token
    Token --> Hash
    Hash --> Vector
    Vector --> Norm
    Norm --> VecItem
    VecItem --> UserIndex

    Query --> Cosine
    UserIndex --> Cosine
    Cosine --> TopK

```

7. Component Dependencies

```

graph TD
    subgraph "Entry Points"
        Main[main.py]
    end

    subgraph "Request/Response Models"
        Schemas[schemas.py]
    end

    subgraph "Core Logic"
        Pipeline[pipeline.py]
        LLM[llm.py]
        Embeddings[embeddings.py]
        VectorStore[vector_store.py]
        Persona[persona.py]
    end

    subgraph "Storage"
        Memory[memory.py]
    end

    subgraph "External"
        Ollama[Ollama API]
        ST[sentence-transformers<br/>optional]
    end

    Main --> Schemas
    Main --> Pipeline
    Main --> Memory
    Pipeline --> LLM
    Pipeline --> Embeddings
    Pipeline --> VectorStore
    Pipeline --> Persona
    Persona --> Embeddings
    LLM --> Ollama
    Embeddings --> ST

    style ST stroke-dasharray: 5 5

```

8. Critique & Scoring Detail

```

flowchart TB
    Text[Generated Text]

    subgraph "Critique Phase"
        BannedList[Global Banned +<br/>User Taboo Phrases]
        Regex[Word-boundary<br/>case-insensitive regex]
        Remove[Remove all matches]
        Normalize[Normalize whitespace<br/>s+→s, \\t→space]
    end

```

```

end

subgraph "Length Check"
    CalcLen[Calculate length]
    Compare[Compare to hint]
    InRange{60-140%<br/>of hint?}
    LengthOK[length_ok = true]
    LengthFail[length_ok = false]
end

subgraph "Style Scoring"
    EmbedText[Embed text]
    GetCentroid[Get persona<br/>centroid]
    CosineSim[cosine(text_emb,<br/>centroid)]
    StyleScore[style_similarity<br/>0.0-1.0]
end

subgraph "Novelty Scoring"
    AllSamples[Get all user<br/>samples]
    MaxSim[Find max similarity<br/>to any sample]
    NoveltyCalc[novelty =<br/>1 - max_sim]
    NoveltyScore[novelty<br/>0.0-1.0]
end

subgraph "Final Score"
    Composite[weighted_score =<br/>style * 0.5 +<br/>novelty * 0.3 +
<br/>constraints * 0.2]
end

Text --> BannedList
BannedList --> Regex
Regex --> Remove
Remove --> Normalize

Normalize --> CalcLen
CalcLen --> Compare
Compare --> InRange
InRange -->|Yes| LengthOK
InRange -->|No| LengthFail

Normalize --> EmbedText
EmbedText --> CosineSim
GetCentroid --> CosineSim
CosineSim --> StyleScore

EmbedText --> MaxSim
AllSamples --> MaxSim
MaxSim --> NoveltyCalc
NoveltyCalc --> NoveltyScore

StyleScore --> Composite
NoveltyScore --> Composite

```

```
LengthOK --> Composite
LengthFail --> Composite
```

9. API Request Flow

```
stateDiagram-v2
    [*] --> HealthCheck: GET /health
    [*] --> CreatePersona: POST /persona
    [*] --> Generate: POST /generate
    [*] --> Stream: POST /generate/stream
    [*] --> Feedback: POST /feedback

    CreatePersona --> PersonaStored
    PersonaStored --> Generate
    PersonaStored --> Stream

    Generate --> VariantsCreated
    Stream --> TokensStreamed

    VariantsCreated --> Feedback
    Feedback --> FeedbackStored

    HealthCheck --> [*]
    FeedbackStored --> [*]
    VariantsCreated --> [*]
    TokensStreamed --> [*]

    note right of Generate
        Requires persona
        Returns scored variants
    end note

    note right of Stream
        SSE format
        No critique/scoring
    end note
```

Usage Notes

These diagrams show:

1. **System Overview:** How components connect
2. **Persona Creation:** Step-by-step user profile building
3. **Generation Pipeline:** Detailed 6-phase generation process
4. **Streaming:** Real-time token generation flow
5. **Data Models:** Entity relationships
6. **Embedding/Retrieval:** Vector search mechanics
7. **Dependencies:** Module interconnections
8. **Critique/Scoring:** Post-processing detail

9. **API Flow:** Request state transitions

Each diagram focuses on a specific aspect to help understand the granular workflow.